

Perceptron Learning Driven Coherence Aware Reuse Prediction for Last-level Caches

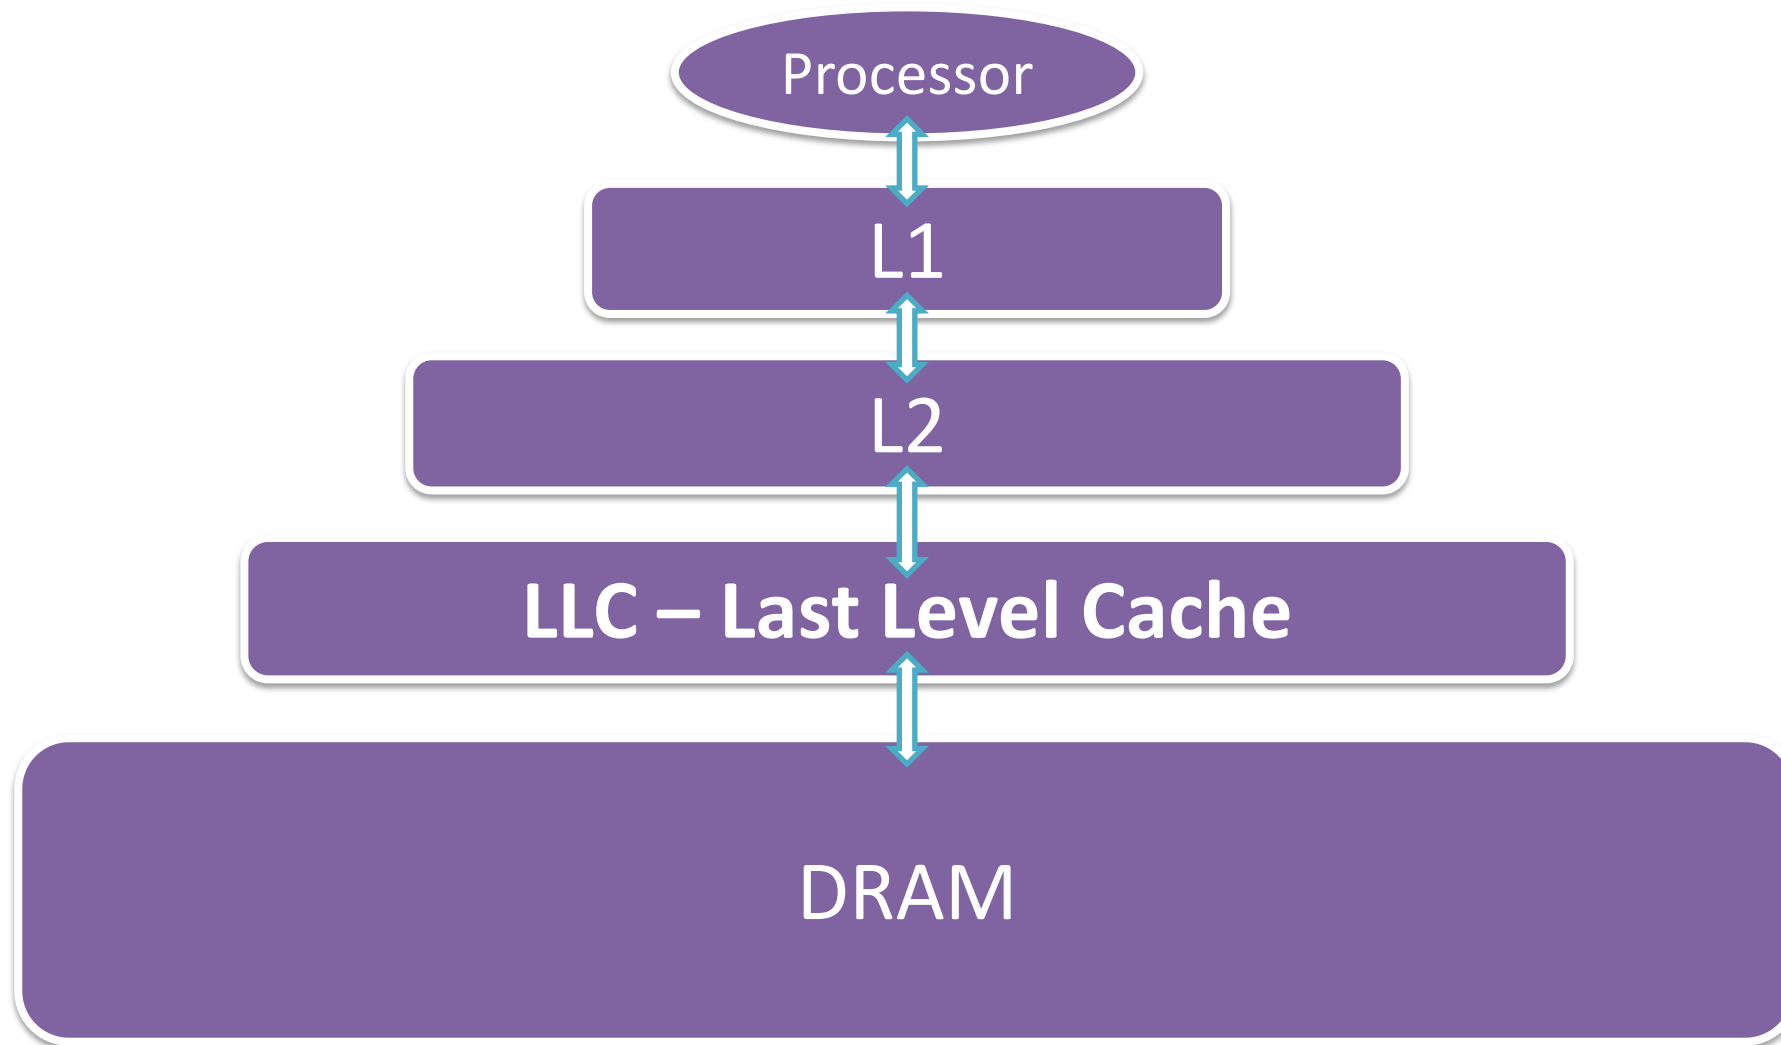
Snehil Verma

Advisor: **Dr. Biswabandan Panda**

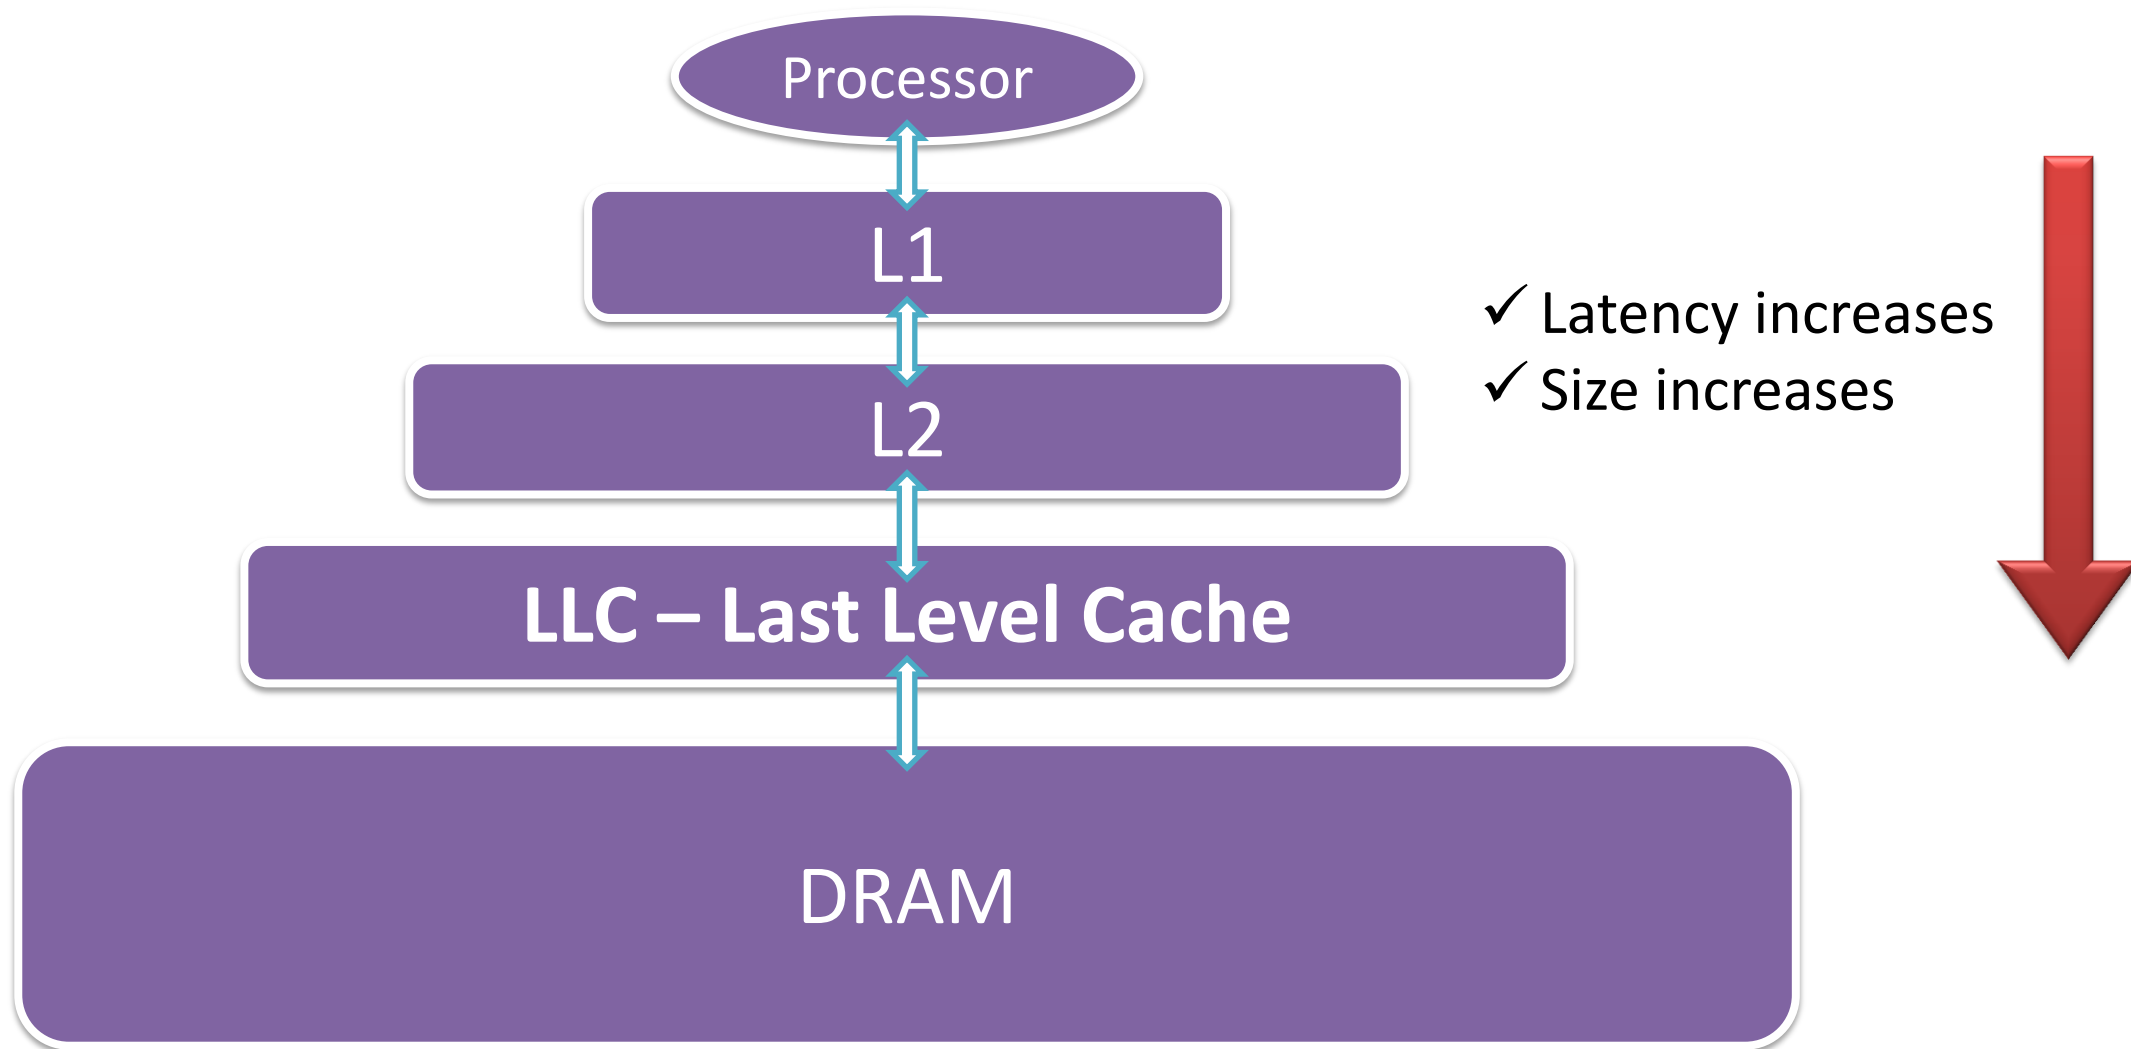


Perceptron Learning Driven
Coherence Aware Reuse Prediction
for **Last-level Caches**

Cache Hierarchy with DRAM



Cache Hierarchy with DRAM



Replacement in Last Level Cache (LLC)

Goal: To minimize off-chip DRAM accesses !!

Replacement in Last Level Cache (LLC)

Goal: To minimize off-chip DRAM accesses !!


- ✓ Preserve important blocks

Replacement in Last Level Cache (LLC)

Goal: To minimize off-chip DRAM accesses !!

- ✓ Preserve important blocks
- ✓ Important blocks → More reused blocks
(enjoy more hits)

Example: Inclusive Cache




Example: Inclusive Cache

Core request



Z

LRU



Example: Inclusive Cache



Core request



Z

Hit

LRU



Example: Inclusive Cache

Core request




Z

5 more requests for Z
(Highly reused block)

Hit

LRU



Example: Inclusive Cache



Core request



5 more requests for Z
(Highly reused block)

Hit

LRU



But no update of LRU status at LLC!



Example: Inclusive Cache

Core request



Miss

LRU




Example: Inclusive Cache

Core request



LRU



Miss




Example: Inclusive Cache

Core request




LRU



Get A

Evict Z




Example: Inclusive Cache


Core request



LRU



Back invalidate Z




Example: Inclusive Cache

Core request



LRU



Get A





Example: Inclusive Cache

Core request



LRU



Highly reused block (Z)
evicted from cache

So, what can be done?

Perceptron Learning Driven
Coherence Aware **Reuse**
Prediction for Last-level Caches

Which blocks have high
reuse ?

Parallel Processing

Parallel Processing

Multi-threaded Applications

Parallel Processing

Multi-threaded Applications

Sharing data b/w multiple Cores

Parallel Processing

Multi-threaded Applications

Sharing data b/w multiple Cores

Shared LLC

Private L1/L2

Shared Blocks in LLC


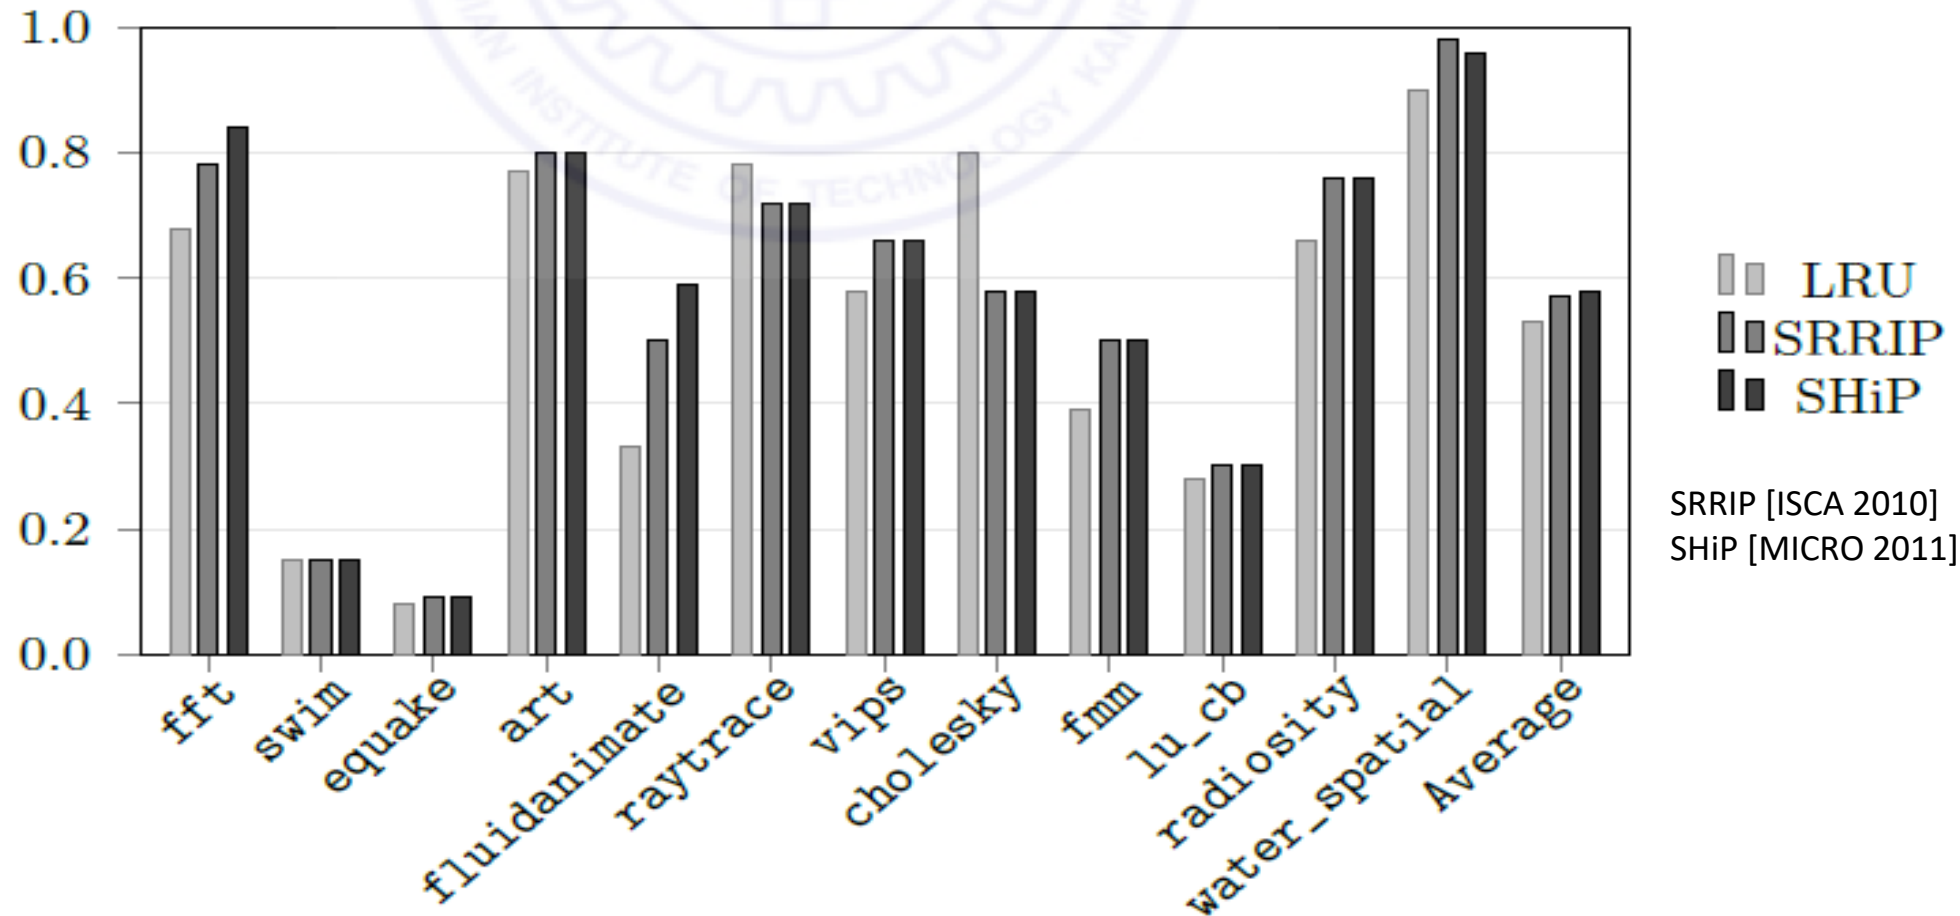


Figure 3.5: Average number of hits enjoyed per block at the LLC with optimal policy

Problem? - Sharing in the Existing LLC Management Policies



SOTA policies:

- **Only 60 %** of cache fills are shared
- They **do not** approximate the optimal sharing behavior well enough

Figure 3.8: Data sharing in baseline policies normalized to optimal sharing


Can we do better ?

How do we predict
reused blocks ?

Perceptron Learning Driven
Coherence *Aware* Reuse Prediction
for Last-level Caches

Perceptron Learning

- **Prediction** of true or false
- Correctness? → **Update weights**




Perceptron Learning Driven
Coherence Aware Reuse
Prediction for Last-level Caches

With parallelism comes
some problems !

Cache Coherence

- **Uniformity** of shared resource data present in multiple local caches
- **Coherency Features:** Sharers ID, Number of sharers, etc.



Main idea

Perceptron based reuse prediction to learn the correlation among the **coherence features** and **reuse** to guide the LLC replacement policy

Five alternatives




- **Bias:** The number of sharers of the requested line is used as bias
- **NumSharers** and **NumSharersHash:** Query the number of sharers
- **OneHot** and **OneHotHash:** One hot encoding of the sharers of a cache block.


Workloads and Replacement Policies

- PARSEC benchmark suite → 8 multi-threaded applications and kernels.
- Combination of two with 4 threads each, using large input data set
- Performance comparison with:
 - SDBP
 - Perceptron Reuse Prediction
 - LRU (Sharers Aware)
 - LRU (no sharers aware)

PARSEC [PACT 2008]
SDBP [MICRO 2010]
Perceptron [MICRO 2016]




Results (4MB LLC) (MPKI)



- *streamcluster* ↔ other benchmarks, except *dedup*, bias → MPKI dropped to at least 25% wrt LRU
- On average, bias → MPKI dropped to 40%, w.r.t LRU
- Other variations of perceptron → lower MPKI w.r.t LRU.

Results (4MB) (Speedup)



- *streamcluster* ↔ other benchmarks, except *dedup*, bias → speedup of at least 40% wrt LRU.
- Bias → geometric mean speedup of 20% over LRU.
- Other variations of perceptron → marginal improvement over LRU.

Conclusion



- We derive five perceptron alternatives and for 4MB & 8MB LLC they all show improvement, on average.
- Specially, perceptron implemented with **bias** outperforms every other replacement policy and shows a major improvement.

Acknowledgement

- This work is the continuation of that initiated at **Texas A&M University**.
- It was performed along with **Jiayi Huang** and **Pritam Majumder**, supervised by **Dr. EJ Kim**.



Thank You!